

LO-PHI: Low-Observable Physical Host Instrumentation for Malware Analysis

Chad Spensky, Hongyi Hu and Kevin Leach

Pietro De Nicolao

Politecnico di Milano

June 13, 2016

Open challenges in Dynamic Malware Analysis

LO-PHI framework implementation

Experiments

Critique

Related and future work

LO-PHI: Low-Observable Physical Host Instrumentation for Malware Analysis

Chad Spensky*†, Hongyi Hu*§ and Kevin Leach*‡

*MIT Lincoln Laboratory *lophi@mit.edu*

†University of California, Santa Barbara *cspensky@cs.ucsb.edu*

§Dropbox *hongyihu@alum.mit.edu*

‡University of Virginia *kjl2y@virginia.edu*

Network and Distributed System Symposium (NDSS) '16
21-24 February 2016, San Diego, CA, USA

- ▶ GitHub repository: <https://github.com/mit-l1/L0-PHI>

Open challenges in Dynamic Malware Analysis

Dynamic Malware Analysis

Idea: **run the malware** in a sandbox (VM, debugger, ...) and use tools to analyze its behavior.

- ▶ We are interested in observing:
 - ▶ Memory accesses
 - ▶ Network activity
 - ▶ Disk activity
- ▶ Observation tools must be placed **outside** the sandbox
- ▶ At a **lower level** w.r.t. the malware
- ▶ In theory, completely *transparent* to the malware
- ▶ *Not so simple...*

Virtualization is like dreaming

- ▶ Unsure if you're living in a dream, or awake?
 - ▶ Look for **artifacts** (i.e. anomalies) in your reality!

Malware can do the same...



Figure 1: Never stopping spinning top: a possible artifact in dreams.

Artifacts and environment-aware malware

Observer effect

Execution of software into a debugger or VM leaves **artifacts**.

- ▶ Artifacts are evidences of an “artificial” environment
- ▶ According to [Garfinkel, 2007], *building a transparent VMM is fundamentally infeasible.*

Malware resistance to dynamic analysis

If the malware is able to detect artifacts, it **can resist to traditional dynamic analysis tools.**

1. Remain inactive (not trigger the payload)
2. Abort or crash its host
3. Disable defenses or tools

Artifacts: some examples

[Chen, 2008], [Garfinkel, 2007] provide a taxonomy of artifacts:

- ▶ **Hardware**

- ▶ Special devices or adapters in VMs
- ▶ Specific manufacturer prefixes on device names

- ▶ **Memory**

- ▶ Hypervisors placing interrupt table in different positions
- ▶ Too little RAM (typical of VMs)

- ▶ **Software**

- ▶ Presence of installed tools on the system
- ▶ `isDebuggerPresent()` Windows API

- ▶ **Behavior**

- ▶ Timing differences (e.g. interception of privileged instructions in VMMs)

Semantic Gap

What is the malware doing?

Need to mine **semantics** from the extracted raw data.

- ▶ From raw data:
 - ▶ Disk read at sector 123
 - ▶ TCP packet ABC
 - ▶ Some bytes at memory address 0xDEADBEEF
- ▶ To concise, high-level event descriptions:
 - ▶ /etc/passwd has been read
 - ▶ A connection to badguy.com has been opened

Forensics to the rescue

LO-PHI uses well-known **forensic tools**, adapted for live analysis:

- ▶ Disk forensics: **Sleuthkit**
- ▶ Memory forensics: **Volatility**

Malware Analysis Framework Evaluation

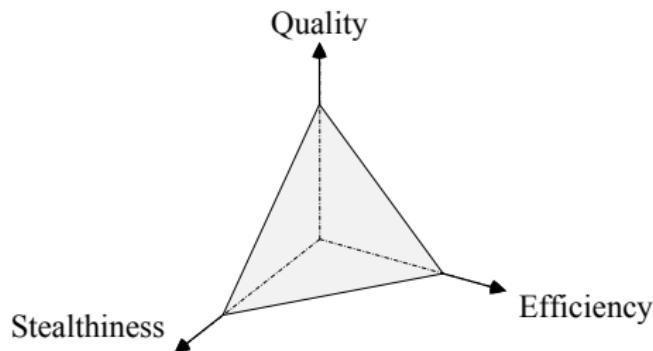


Figure 2: From [Kirat, 2011]

Tradeoff between:

1. *Low-artifact, semantically poor* tools (Virtual Machine Introspection)
2. *High-artifact, semantically-rich* frameworks (debuggers)

LO-PHI framework implementation

Goals

- ▶ **No virtualization:** run malware on bare metal machines!
 - ▶ Physical sensors and actuators
- ▶ Bridging the **semantic gap**
 - ▶ Physical sensors collect raw data
- ▶ **Automated restore** to pre-infection state
- ▶ **Stealthiness:** very few, undetectable artifacts
- ▶ **Extendability:** support new OSs and filesystems

Threat model

Assumptions on our model of malware: they are **limitations** of the approach.

1. Malicious modifications evident either in memory or on disk
2. No infection delivered to hardware
3. Instrumentation is in place before malware is executed
 - ▶ Malware cannot analyze the system without LO-PHI in place
 - ▶ Harder to compare and detect artifacts: **no baseline**

Sensors and actuators (1)

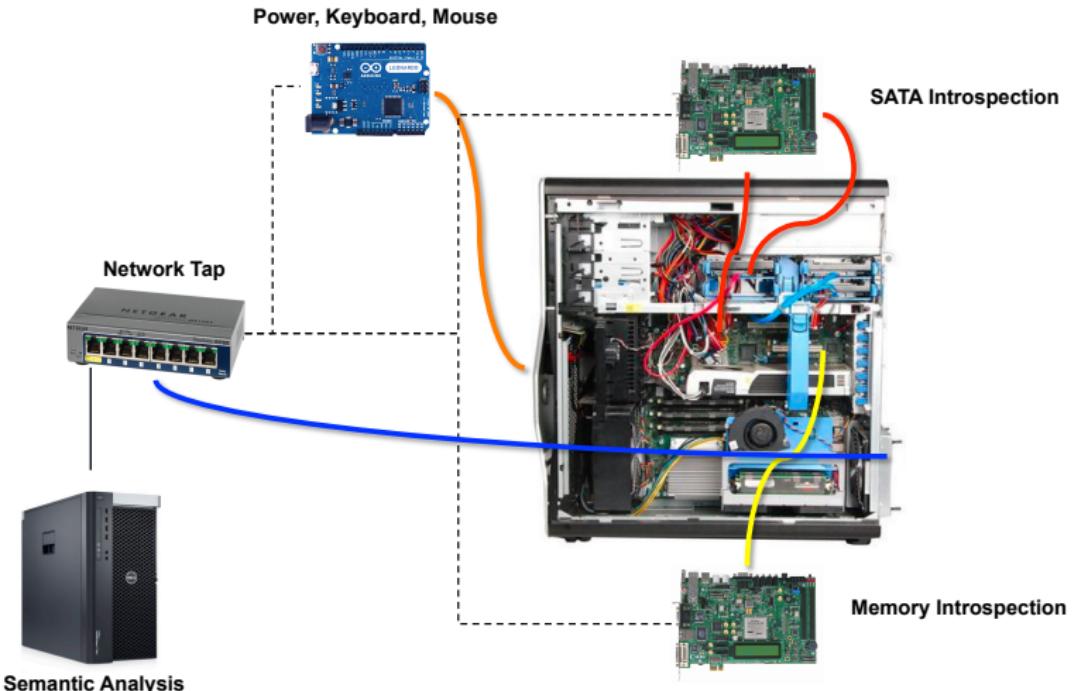


Figure 3: Hardware instrumentation to inspect a bare-metal machine.

Sensors and actuators (2)

Sensors

- ▶ **Memory.** Xilinx ML507 board connected to PCIe, reads and writes arbitrary memory locations via DMA.
- ▶ **Disk.** ML507 board intercepting all the traffic over SATA interface. Sends SATA frames via Gigabit Ethernet and UDP.
 - ▶ Completely passive...
 - ▶ except when SATA data rate exceeds Ethernet bandwidth: throttling of frames.
- ▶ **Network interface.** Mentioned in paper, but the technology used is unclear.

Actuators

An Arduino Leonardo emulates **USB keyboard and mouse**.

Infrastructure

Restoring physical machines

- ▶ We cannot simply “restore a snapshot” like in VMs
- ▶ **Preboot Execute Environment (PXE)** with **CloneZilla**
 - ▶ Allows to restore the disk to a previously saved state
 - ▶ No interaction with the OS

Scalable infrastructure

- ▶ Job submission system: jobs are sent to a **scheduler**
- ▶ The scheduler executes the routine on an appropriate machine
- ▶ **Python interface** to control the machine and run malware and analysis

Bridging the semantic gap via forensic analysis

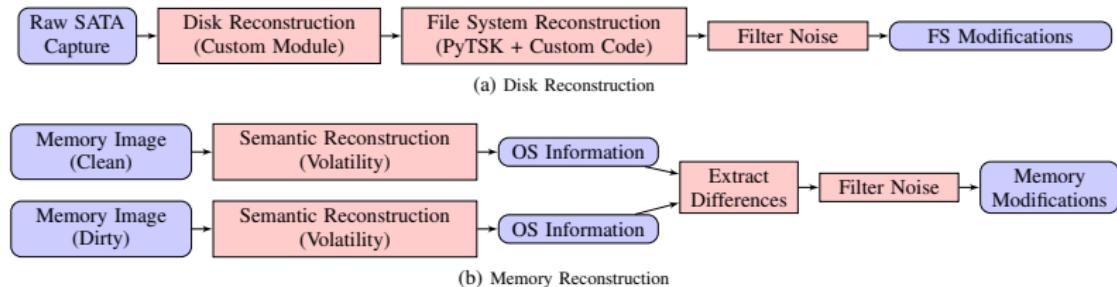


Figure 4: Binary analysis workflow. Rounded nodes represent data and rectangles represent data manipulation.

- ▶ **Background noise** was removed by analyzing non-malicious software.

Example of output of the analysis toolbox

Offset	Name	PID	PPID
0x86292438	AcroRd32.exe	1340	1048
0x86458818	AcroRd32.exe	1048	1008
0x86282be0	AdobeARM.exe	1480	1048
0x864562a0	\$\$_rk_sketchy_server.exe	1044	1008

Figure 5: New processes.

PID	Port	Protocol	Address
1048	1038	UDP	127.0.0.1
1044	21	TCP	0.0.0.0

Figure 6: New sockets.

Created Filename
.../lophi/\$\$ rk_sketchy_server.exe
.../lophi/hookssdt.sys
.../lophi/sample_0742475e94904c41de1397af5c53dff8e.exe

Figure 7: Disk event log.

Experiments

Experiment on evasive malware

Malware previously labeled as “evasive” was executed on Windows 7 and analyzed using LO-PHI.

- ▶ Many samples clearly exhibited **typical malware behavior**.
- ▶ Some samples failed because of no network connection, or wrong Windows version.

[...] we feel that our findings are more than sufficient to showcase LO-PHI’s ability to analyze evasive malware, without being subverted, and subsequently produce high-fidelity results for further analysis.

Timing

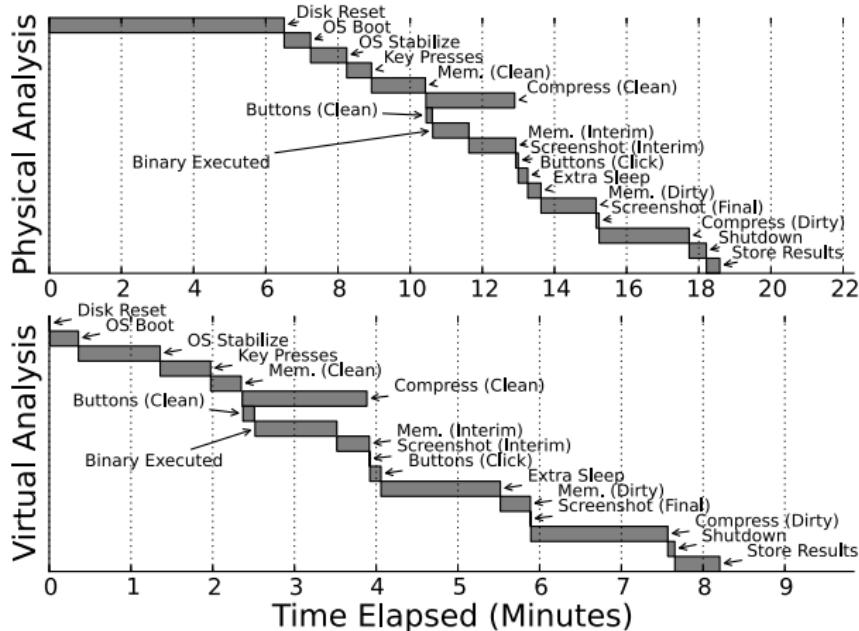


Figure 8: Time spent in each step of binary analysis. Both environments were booting a 10 GB Windows 7 (64-bit) hibernate image and were running on a system with 1 GB of volatile memory.

Critique

Known limitations

- ▶ Newer chipsets use IOMMUs, **disabling DMA** from peripherals
 - ▶ Current memory acquisition technique will become unusable
- ▶ **Smearing:** the memory can change *during* the acquisition
 - ▶ Inconsistent states
 - ▶ Faster polling rates can help
- ▶ **Filesystem caching:** some data will not pass through SATA interface
 - ▶ Malware could write a file to disk cache, execute and delete it before the cached is flushed to disk.
 - ▶ However, the effects would be visible in memory.

Issues with the technique and the experiments

- ▶ The malware is left to run only 3 minutes.
 - ▶ Many malwares need much more time to fully uncover their effects (e.g. ransomware).
- ▶ **No memory polling** during the execution of the malware
 - ▶ Only snapshot before and after the execution
 - ▶ Temporary data used by the malware is never seen
- ▶ Assumption: **malware does not modify BIOS or firmware.**
 - ▶ But if it does, the physical machine could not be recoverable.
 - ▶ Costly!
- ▶ **No Internet access:** the authors always run the malware on disconnected machines.
 - ▶ Most malware becomes useless without Command&Control infrastructure.
 - ▶ Network access could expose further, unseen, artifacts.

Methodological issues

- ▶ The article claims that the artifacts from LO-PHI are unusable by malware because there's **no baseline** (i.e. the malware cannot see the machine before the installation of LO-PHI).
 - ▶ This can also be true for traditional approaches.
- ▶ **No statistical test** used to discover whether difference in disk/memory throughput is statistically significant (presence of artifacts).
 - ▶ Very simple and standard procedure, should really be done in a scientific paper.
- ▶ **Network analysis** technique is not described and unclear.
 - ▶ *We exclude the network trace analysis from much of our discussion since it is a well-known technique and not the focus of our work.*

Memory throughput with and without LO-PHI

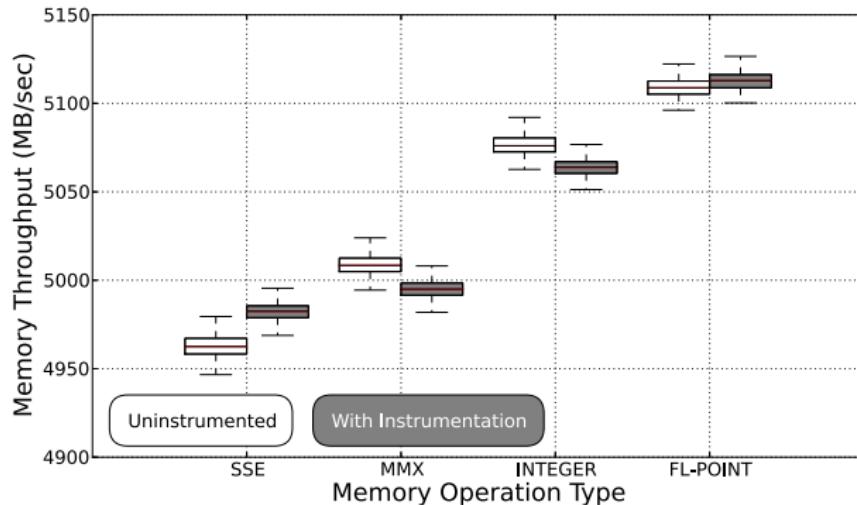


Figure 9: Average memory throughput comparison as reported by RAMSpeed, with and without instrumentation. Deviation from uninstrumented trial is only 0.4% in worst case. **No statistical test used.**

Another point...

- ▶ **Virtualization is increasingly used** in production contexts.
 - ▶ Just think about **cloud computing!**
- ▶ Attackers will want to target those VMs [Garfinkel, 2007]
- ▶ Malware that deactivates in VMs will be less common

Related and future work

Related work

Traditional dynamic analysis

Many dynamic malware analysis tools rely on virtualization: Ether, BitBlaze, Anubis, V2E, HyperDbg, SPIDER.

- ▶ We already saw the limitations of VM approaches: **artifacts**

Bare-metal dynamic analysis

BareBox [Kirat, 2011]: malware analysis framework based on a bare-metal machine without virtualization or emulations

- ▶ Only targets user-mode malware
- ▶ Only disk analysis (no memory tools)

Future evolutions

Automated, repeated analyses

- ▶ **Disk restore** phase is the lengthiest (> 6 min)
- ▶ *The resetting and boot process could be decreased significantly by writing a custom PXE loader, or completely mitigated by implementing copy-on-write into our FPGA.*
- ▶ *While snapshots are trivial with virtual machines, it is still an open problem for physical machines.*

Extensions

- ▶ Analyze **transient behavior** of binary
 - ▶ Continuous memory polling
 - ▶ Need to deal with DMA artifacts
- ▶ Cover **malware that infects hardware** (BIOS, firmware)

References

-  Chen X., Andersen J., Morley M., Bailey M., Nazario J.
Towards an understanding of anti-virtualization and anti-debugging behavior in modern malware
In Proceedings of the International Conference on Dependable Systems and Networks (2008)
doi: 10.1109/DSN.2008.4630086
-  Kirat D., Vigna G., Kruegel C.
BareBox: Efficient Malware Analysis on Bare-metal
Proceedings of the 27th Annual Computer Security Applications Conference (2001)
doi: 10.1145/2076732.2076790
-  Garfinkel T., Adams K., Warfield A., Franklin J.
Compatibility is Not Transparency: VMM Detection Myths and Realities
Proceedings of the 11th USENIX Workshop on Hot Topics in Operating Systems (2007)